

*"Mathematics is the art of giving the
same name to different things."
- Henri Poincaré*

Extra Project

Randomized SVD

Abror Shopulatov Mohammed Ibrahim Awad Imran Turganov

Mohamed bin Zayed University of Artificial Intelligence

November, 2025

Recall

During the course, we saw:

- ▶ Vectors and matrices
- ▶ Special matrices: rotations, reflections, projections
- ▶ Eigenvalues and eigenvectors — directions that matrices simply scale
- ▶ That lead us to matrix decompositions and SVD

Remark: Singular Value Decomposition (SVD)

Any matrix $A \in \mathbb{R}^{m \times n}$ can be factored as:

$$A = U\Sigma V^\top$$

where:

- ▶ $U \in \mathbb{R}^{m \times m}$ and $V \in \mathbb{R}^{n \times n}$ are orthogonal matrices
- ▶ $\Sigma \in \mathbb{R}^{m \times n}$ is diagonal with $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r \geq 0$

SVD usage

SVD is everywhere in data science:

- ▶ Image and video compression, Principal Component Analysis (PCA)
- ▶ Recommender systems, Noise reduction and signal processing

Example (Image Compression with SVD)



Original Image
(full rank)



Rank-100 Approximation
Stores only 9.77% of original!

The Challenge

SVD works beautifully... but what happens we go higher dimentions?

Lemma: Classical SVD Complexity

Computing the full SVD of an $m \times n$ matrix requires:

$$\mathcal{O}(\min\{mn^2, m^2n\})$$

Matrix Size	Operations	Time
$1,000 \times 1,000$	$\sim 10^9$	seconds
$10,000 \times 10,000$	$\sim 10^{12}$	hours
$100,000 \times 100,000$	$\sim 10^{15}$	infeasible

The Problem

Key Observation

Computing SVD of **large matrices** is computationally **expensive!**

What we could do:

1. Work on faster *computers*
2. Look for smarter *algorithms*
3. Give up!

The real problem:

- ▶ **Classical SVD:** Computes *all* singular vectors and values
- ▶ **Our goal:** Capture only the **top- k** singular vectors

The Randomized Idea - Intuition

Key Insight: Random sampling preserves geometric structure with high probability

Example: Random Projection

Draw a random matrix $\Omega \in \mathbb{R}^{n \times k}$ with Gaussian entries, then:

$$Y = A\Omega \in \mathbb{R}^{m \times k}$$

What does Y capture?

- ▶ Each column of Y "probes" the range of A
- ▶ Large singular values of A dominate the response
- ▶ Small singular values contribute negligibly
- ▶ Result: column space of $Y \approx$ top- k subspace of A

Randomized SVD Algorithm

Definition: Algorithm: Randomized SVD

Input: Matrix $A \in \mathbb{R}^{m \times n}$, target rank k where $k \ll \min\{m, n\}$


Steps:

1. **Random sketch:** Draw $\Omega \in \mathbb{R}^{n \times k}$ with i.i.d. Gaussian entries $\mathcal{N}(0, 1)$
Compute $Y = A\Omega$ $[O(mnk)]$
2. **Orthogonalize:** $Q = \text{orth}(Y)$ via QR factorization $[O(mk^2)]$
3. **Project:** $B = Q^\top A$ $[O(mnk)]$
4. **Small SVD:** Compute $B = \hat{U}_B \Sigma V^\top$ $[O(nk^2)]$
5. **Reconstruct:** $U = Q \hat{U}_B$ $[O(mk^2)]$

Output: $U \in \mathbb{R}^{m \times k}$, $\Sigma \in \mathbb{R}^{k \times k}$, $V \in \mathbb{R}^{n \times k}$

Total complexity: $O(mnk)$ vs. $O(\min\{mn^2, m^2n\})$ for classical SVD

Visualization



Error Bounds

Lemma: Halko-Martinsson-Tropp (for the brave it can be omitted)

Let $A \in \mathbb{R}^{m \times n}$ with singular values $\sigma_1 \geq \dots \geq \sigma_n$. Let $\Omega \in \mathbb{R}^{n \times (k+p)}$ have i.i.d. Gaussian entries, and $Q = \text{orth}(A\Omega)$. Then:

Frobenius norm:

$$\mathbb{E}[\|A - QQ^\top A\|_F] \leq \left(1 + \frac{k}{p-1}\right)^{1/2} \left(\sum_{i=k+1}^n \sigma_i^2\right)^{1/2}$$

Spectral norm:



$$\mathbb{E}[\|A - QQ^\top A\|_2] \leq \left(1 + \frac{k}{p-1}\right)^{1/2} \sigma_{k+1}$$

Remark:

With oversampling $p = 10$, error is within $\sim 3\times$ of optimal rank- k approximation!



Numerical Results

Experimental setup: Random Gaussian matrices, $n \in \{10^2, 10^3, 10^4, 10^5, 10^6\}$, rank $k = 200$



Classical SVD

Cubic scaling makes it infeasible beyond 10^4



Randomized SVD

Scales to $10^6 \times 10^6$ with $q \in \{0, 1, 2, 3, 4\}$

Key observations:

- ▶ Classical SVD: 130+ seconds at $10^4 \times 10^4$
- ▶ rSVD: <1 second at $10^6 \times 10^6$ (with $q = 1$: 318 ms)
- ▶ Power iteration ($q \geq 2$) closes accuracy gap to machine precision

Conclusion

What we've shown:

- ▶ Classical SVD is powerful but computationally expensive: $\mathcal{O}(\min\{mn^2, m^2n\})$
- ▶ **Randomized SVD** achieves $\mathcal{O}(mn \log(k))$ complexity through random sampling
- ▶ **Rigorous error guarantees:** Near-optimal with small oversampling
- ▶ **Practical performance:** 10-100 \times speedup with negligible accuracy loss

Thank you!

Full paper and code: <https://github.com/IMRUNya/rSVD>